

# Challenges + job category

## Overview

Secondary analysis of survey Q9: “How frequently have you encountered the following challenges while working on open-source projects?”

In this script, I am considering challenges in light of job category.

### Import packages and utilities

```
project_root <- here::here() # requires that you be somewhere in the  
# project directory (not above it)  
# packages  
suppressMessages(source(file.path(project_root, "scripts/packages.R")))  
# functions and objects used across scripts  
suppressMessages(source(file.path(project_root, "scripts/utils.R")))
```

### Load data

```
challenges <- load_qualtrics_data("clean_data/challenges_Q9.tsv")  
other_quant <- load_qualtrics_data("clean_data/other_quant.tsv")
```

### Wrangle data

```
challenges_and_job <- challenges  
challenges_and_job$job_category <- other_quant$job_category  
  
head(challenges_and_job)
```

```

Coding time Documentation time Managing issues Attracting users Recognition
1 Always Always Always Always Always Always
2 Frequently Occasionally Occasionally Occasionally Occasionally
3 Frequently Always Occasionally Always Occasionally
4 Always Always Frequently Occasionally Frequently
5 Always Always Rarely Occasionally Frequently
6

Hiring Security Finding peers Finding mentors Education time
1 Always Always Always Always Always
2 Rarely Frequently Occasionally Frequently Frequently
3 Frequently Frequently Occasionally Occasionally Rarely
4 Always Occasionally Rarely Rarely Frequently
5 Never Never Never Never Always
6

Educational resources Legal Finding funding Securing funding
1 Always Always Always Always
2 Frequently Frequently Frequently Occasionally
3 Rarely Always Always Always
4 Rarely Occasionally Frequently Frequently
5 Occasionally Occasionally Rarely Always
6

job_category
1 Faculty
2 Post-Doc
3 Other research staff
4 Faculty
5 Faculty
6 Other research staff

```

Remove rows that contain any empty entries.

```
nrow(challenges_and_job)
```

```
[1] 332
```

```
challenges_and_job <- exclude_empty_rows(challenges_and_job, strict = TRUE) # from scripts/u
nrow(challenges_and_job)
```

```
[1] 233
```

For visual clarity in our plots, let's combine postdocs and other staff researchers, as well as undergrads and grad students.

```

challenges_and_job <- challenges_and_job %>%
  mutate(
    job_category = recode(
      job_category,
      "Post-Doc" = "Postdocs and Staff Researchers",
      "Other research staff" = "Postdocs and Staff Researchers"
    )
  )

challenges_and_job <- challenges_and_job %>%
  mutate(
    job_category = recode(
      job_category,
      "Grad Student" = "Students",
      "Undergraduate" = "Students"
    )
  )

challenges_and_job$participantID <- row.names(challenges_and_job)

head(challenges_and_job)

```

	Coding time	Documentation time	Managing issues	Attracting users	Recognition
1	Always	Always	Always	Always	Always
2	Frequently	Occasionally	Occasionally	Occasionally	Occasionally
3	Frequently	Always	Occasionally	Always	Occasionally
4	Always	Always	Frequently	Occasionally	Frequently
5	Always	Always	Rarely	Occasionally	Frequently
7	Frequently	Frequently	Frequently	Frequently	Frequently
	Hiring	Security	Finding peers	Finding mentors	Education time
1	Always	Always	Always	Always	Always
2	Rarely	Frequently	Occasionally	Frequently	Frequently
3	Frequently	Frequently	Occasionally	Occasionally	Rarely
4	Always	Occasionally	Rarely	Rarely	Frequently
5	Never	Never	Never	Never	Always
7	Always	Never	Never	Never	Frequently
	Educational resources		Legal	Finding funding	Securing funding
1		Always	Always	Always	Always
2		Frequently	Frequently	Frequently	Occasionally

	Rarely	Always	Always	Always
	Rarely	Occasionally	Frequently	Frequently
	Occasionally	Occasionally	Rarely	Always
	Never	Always	Always	Always
	job_category participantID			
1	Faculty		1	
2	Postdocs and Staff Researchers		2	
3	Postdocs and Staff Researchers		3	
4		Faculty	4	
5		Faculty	5	
7		Faculty	7	

Let's reshape the data from wide to long format for easier counting and plotting.

```
long_data <- challenges_and_job %>%
  pivot_longer(
    cols = -c(participantID, job_category),
    names_to = "challenge",
    values_to = "challenge_level"
  )

long_data

# A tibble: 3,262 x 4
  job_category participantID challenge      challenge_level
  <chr>        <chr>       <chr>        <chr>
  1 Faculty      1            Coding time   Always
  2 Faculty      1            Documentation time Always
  3 Faculty      1            Managing issues Always
  4 Faculty      1            Attracting users Always
  5 Faculty      1            Recognition    Always
  6 Faculty      1            Hiring         Always
  7 Faculty      1            Security       Always
  8 Faculty      1            Finding peers   Always
  9 Faculty      1            Finding mentors Always
 10 Faculty     1            Education time Always
# i 3,252 more rows
```

## Top 5 challenges per group, by “points”

```
long_data <- long_data %>%
  mutate(
    challenge_score = recode(
      challenge_level,
      "Never"          = 0L,
      "Non-applicable" = 0L,
      "Rarely"         = 1L,
      "Occasionally"   = 2L,
      "Frequently"     = 3L,
      "Always"         = 4L
    )
  )
# Using interger literals 0L, 1L, etc., ensures that
# the new column will be integers, not doubles.

long_data
```

```
# A tibble: 3,262 x 5
  job_category participantID challenge      challenge_level challenge_score
  <chr>           <chr>       <chr>           <chr>                  <int>
1 Faculty          1            Coding time    Always                   4
2 Faculty          1            Documentation time Always                   4
3 Faculty          1            Managing issues Always                   4
4 Faculty          1            Attracting users Always                  4
5 Faculty          1            Recognition    Always                   4
6 Faculty          1            Hiring          Always                   4
7 Faculty          1            Security        Always                   4
8 Faculty          1            Finding peers   Always                   4
9 Faculty          1            Finding mentors Always                   4
10 Faculty         1            Education time Always                  4
# i 3,252 more rows
```

Let's just inspect all the basic stats for all the challenges for each group.

```
# Helper to compute the (numeric) mode
get_mode <- function(x) {
  ux <- unique(x)
  ux[which.max(tabulate(match(x, ux)))]
```

```

}

get_summary_df <- function(job_str, df) {
  res <- df %>%
    filter(job_category == job_str) %>%
    group_by(challenge) %>%
    summarise(
      total = sum(challenge_score),
      mean = mean(challenge_score, na.rm = TRUE),
      median = median(challenge_score),
      mode = get_mode(challenge_score),
      st_dev = sd(challenge_score, na.rm = TRUE)
    ) %>%
    ungroup() %>%
    arrange(desc(total))
  return(res)
}

jobs_ordered <- c(
  "Faculty",
  "Postdocs and Staff Researchers",
  "Students",
  "Non-research Staff"
)

summary_tables <- lapply(jobs_ordered, function(j) get_summary_df(j, long_data))
names(summary_tables) <- jobs_ordered
summary_tables

$Faculty
# A tibble: 14 x 6
  challenge      total   mean  median   mode  st Dev
  <chr>        <int> <dbl> <int> <int> <dbl>
1 Documentation  176    2.98     3     4   1.22
2 Coding         170    2.88     3     4   1.31
3 Securing funding  158    2.68     3     4   1.58
4 Finding funding  149    2.53     3     4   1.57
5 Education time  142    2.41     2     4   1.35
6 Managing issues  115    1.95     2     3   1.41
7 Hiring          114    1.93     2     0   1.65
8 Attracting users  113    1.92     2     2   1.34
9 Recognition     108    1.83     2     0   1.52

```

```

10 Educational resources      84  1.42      1      0  1.29
11 Legal                      84  1.42      1      0  1.32
12 Finding mentors            76  1.29      1      0  1.40
13 Finding peers              76  1.29      1      0  1.19
14 Security                   76  1.29      1      0  1.35

$`Postdocs and Staff Researchers`
# A tibble: 14 x 6
  challenge      total   mean median mode st_dev
  <chr>        <int> <dbl>  <int> <int> <dbl>
1 Documentation time     174  3.16    3      3  0.788
2 Coding time             145  2.64    3      3  1.16
3 Education time          133  2.42    3      3  1.30
4 Securing funding         120  2.18    3      4  1.72
5 Finding funding           118  2.15    2      4  1.60
6 Attracting users          116  2.11    2      3  1.49
7 Managing issues           113  2.05    2      2  1.30
8 Recognition                 88  1.6      2      2  1.24
9 Educational resources       85  1.55    1      1  1.15
10 Legal                      84  1.53    1      1  1.17
11 Finding mentors            73  1.33    1      0  1.28
12 Hiring                      68  1.24    0      0  1.59
13 Security                   62  1.13    1      0  1.19
14 Finding peers              58  1.05    1      0  0.970

$Students
# A tibble: 14 x 6
  challenge      total   mean median mode st_dev
  <chr>        <int> <dbl>  <int> <int> <dbl>
1 Coding time             85  2.58    3      3  1.12
2 Documentation time       84  2.55    3      3  1.37
3 Education time            75  2.27    2      3  1.21
4 Attracting users           59  1.79    2      0  1.54
5 Educational resources       59  1.79    2      1  1.27
6 Managing issues             57  1.73    2      2  1.35
7 Finding mentors             53  1.61    2      0  1.25
8 Finding funding              52  1.58    1      0  1.64
9 Securing funding             49  1.48    1      0  1.64
10 Recognition                  46  1.39    1      0  1.32
11 Legal                      42  1.27    1      0  1.35
12 Finding peers                 41  1.24    1      0  1.35
13 Security                     26  0.788   0      0  1.34
14 Hiring                      18  0.545   0      0  1.23

```

```
$`Non-research Staff`  
# A tibble: 14 x 6  
  challenge      total   mean median mode st_dev  
  <chr>        <int> <dbl>  <dbl> <int> <dbl>  
1 Documentation time     252  2.93    3     3  0.968  
2 Coding time            206  2.40    3     3  1.26  
3 Education time         189  2.20    2     3  1.21  
4 Managing issues        166  1.93    2     2  1.18  
5 Attracting users       154  1.79    2     0  1.48  
6 Security               143  1.66    2     2  1.28  
7 Educational resources   141  1.64    2     1  1.11  
8 Legal                  123  1.43    1     2  1.20  
9 Finding mentors         121  1.41    1     0  1.30  
10 Finding funding        113  1.31    0     0  1.63  
11 Securing funding       111  1.29    0     0  1.65  
12 Finding peers          92   1.07    1     0  1.09  
13 Recognition            92   1.07    1     0  1.23  
14 Hiring                 91   1.06    0     0  1.35
```

Let's plot the top 5. First, a little data wrangling.

```
all_tbl <- bind_rows(summary_tables, .id = "job_category")  
  
top5_by_job <- all_tbl %>%  
  group_by(job_category) %>%  
  slice_max(mean, n = 5, with_ties = FALSE) %>%  
  ungroup() %>%  
  select(job_category, challenge, mean) %>%  
  mutate(job_category = factor(job_category, levels = jobs_ordered))  
  
# Reorder factor levels for visual clarity  
ordered_challenges_top5_by_points <- c(  
  "Documentation time",  
  "Coding time",  
  "Education time",  
  "Educational resources",  
  "Securing funding",  
  "Finding funding",  
  "Managing issues",  
  "Attracting users"  
)
```

```
top5_by_job$challenge <- factor(  
  top5_by_job$challenge,  
  levels = ordered_challenges_top5_by_points  
)
```

Let's add a whitespace in this long job category name

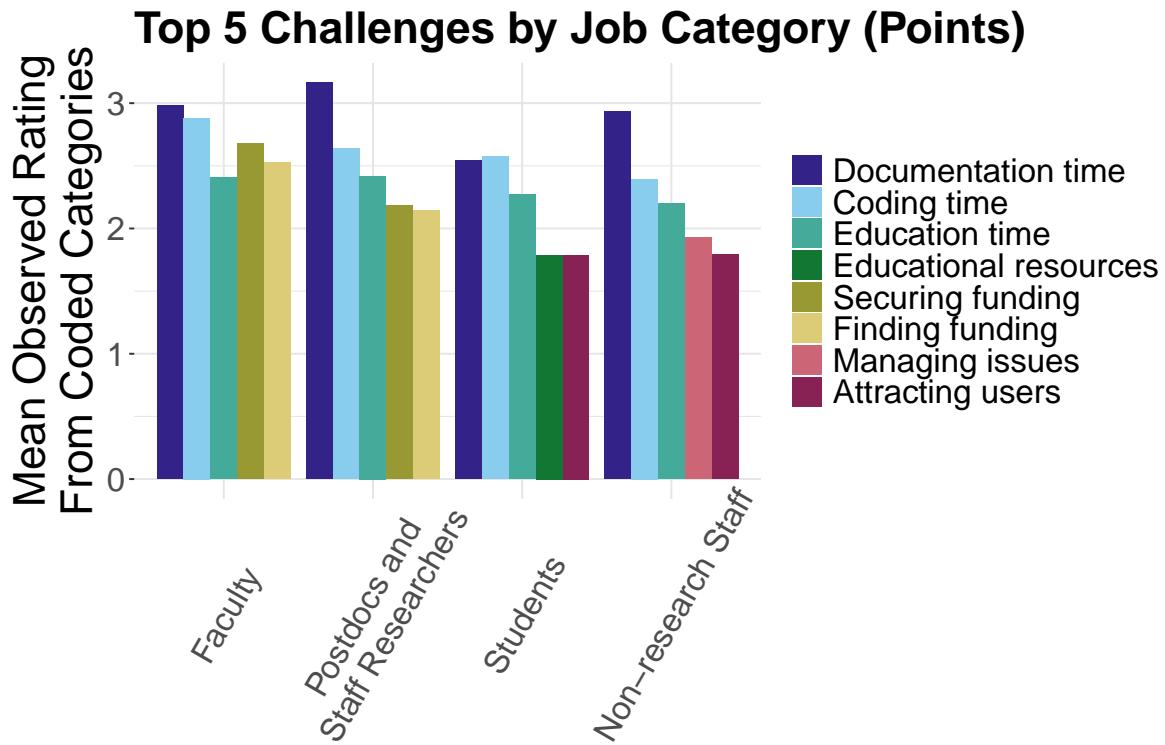
```
top5_by_job <- top5_by_job %>%  
  mutate(  
    job_category = recode(  
      job_category,  
      "Postdocs and Staff Researchers" = "Postdocs and\nStaff Researchers",  
    )  
  )
```

Let's hard-code a color palette that is tailored to these data. This will be useful in the next section, when we plot almost the same set of challenges, and we'll want the challenges to correspond to the same colors in the legend.

```
# I'm just including the names here for my own reference,  
# but they're not actually used in the code.  
chall_colors <- list(  
  # modified from https://sronpersonalpages.nl/~pault/  
  "Documentation time" = "#332288",  
  "Coding time" = "#88CCEE",  
  "Education time" = "#44AA99",  
  "Educational resources" = "#117733",  
  "Securing funding" = "#999933",  
  "Finding funding" = "#DDCC77",  
  "Managing issues" = "#CC6677",  
  "Attracting users" = "#882255"  
)
```

```
top5_by_points_plot <- ggplot(  
  top5_by_job,  
  aes(  
    x = job_category,  
    y = mean,  
    fill = challenge  
)  
) +
```

```
geom_col(position = position_dodge()) +
scale_fill_manual(values = chall_colors) +
labs(
  x = "Job Category",
  y = "Mean Observed Rating\nFrom Coded Categories",
  fill = "Challenge",
  title = "Top 5 Challenges by Job Category (Points)"
) +
theme(
  axis.title.x = element_blank(),
  axis.title.y = element_text(size = 24),
  axis.text.x = element_text(angle = 60, vjust = 0.6, size = 18),
  axis.text.y = element_text(size = 18),
  axis.ticks.x = element_blank(),
  legend.title = element_blank(),
  legend.text = element_text(size = 18),
  panel.background = element_blank(),
  panel.grid = element_line(linetype = "solid", color = "gray90"),
  plot.title = element_text(hjust = 0, size = 24, face = "bold"),
  plot.margin = unit(c(0.3, 0.3, 0.3, 0.3), "cm")
)
top5_by_points_plot
```



```
save_plot("top5_challenges_by_job_points.tiff", 12, 10, p=top5_by_points_plot)
```

### Top 5 challenges per group, by proportion frequently or always

As another way of confirming/exploring these trends, let's look at the proportion of each group who said "frequently" or "always".

```
# Calculate proportion of TRUEs by taking the mean of a logical vector,
# created by %in%.
proportions <- long_data %>%
  group_by(job_category, challenge) %>%
  summarize(proportion = mean(challenge_level %in% c("Frequently", "Always"))) %>%
  ungroup()
```

``summarise()` has grouped output by 'job_category'. You can override using the `groups` argument.`

```
proportions
```

```
# A tibble: 56 x 3
  job_category challenge      proportion
  <chr>        <chr>          <dbl>
1 Faculty      Attracting users 0.356
2 Faculty      Coding time    0.712
3 Faculty      Documentation time 0.763
4 Faculty      Education time 0.492
5 Faculty      Educational resources 0.186
6 Faculty      Finding funding 0.627
7 Faculty      Finding mentors 0.220
8 Faculty      Finding peers   0.169
9 Faculty      Hiring          0.475
10 Faculty     Legal           0.169
# i 46 more rows
```

```
top5_by_prop <- proportions %>%
  group_by(job_category) %>%
  slice_max(order_by = proportion, n = 5)
```

```
# Filter to include only challenges present in the top5 dataframe
filtered_props <- proportions %>%
  semi_join(top5_by_prop, by = c("job_category", "challenge"))
```

Let's inspect the challenges that made the cut. Are they the same as the challenges from calculating the top 5 the other way (by points)?

```
ordered_challenges_top5_by_points
```

```
[1] "Documentation time"      "Coding time"            "Education time"
[4] "Educational resources"   "Securing funding"       "Finding funding"
[7] "Managing issues"         "Attracting users"
```

```
unique(filtered_props$challenge)
```

```
[1] "Coding time"           "Documentation time" "Education time"
[4] "Finding funding"        "Securing funding"    "Attracting users"
[7] "Managing issues"
```

They are not the same. In this case, educational resources didn't make the cut. We can still use the same order of factor levels.

```
# Reorder factor levels
filtered_props$challenge <- factor(
  filtered_props$challenge,
  levels = ordered_challenges_top5_by_points
)

filtered_props$job_category <- factor(
  filtered_props$job_category,
  levels = jobs_ordered
)
```

Let's add a whitespace in this long job category name

```
filtered_props <- filtered_props %>%
  mutate(
    job_category = recode(
      job_category,
      "Postdocs and Staff Researchers" = "Postdocs and\nStaff Researchers",
    )
  )
```

Custom color palette, which doesn't contain “Educational resources”

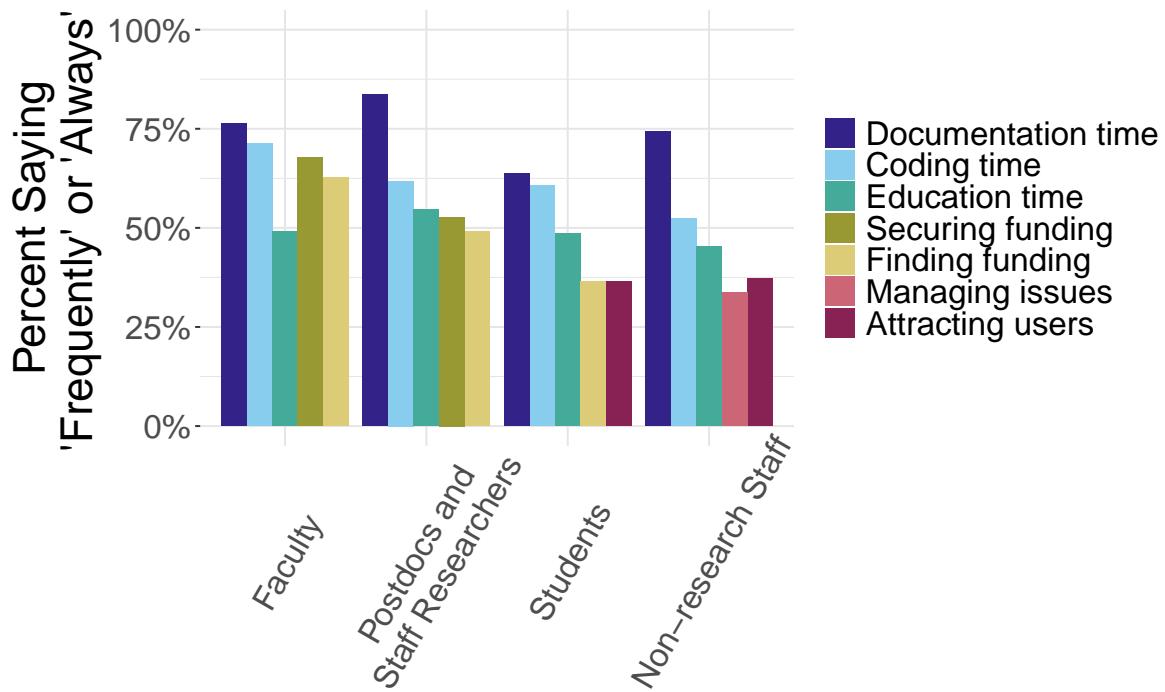
```
# I'm just including the names here for my own reference,
# but they're not actually used in the code.
chall_colors2 <- list(
  # modified from https://sronpersonalpages.nl/~pault/
  "Documentation time" = "#332288",
  "Coding time" = "#88CCEE",
  "Education time" = "#44AA99",
  # "Educational resources" = "#117733",
  "Securing funding" = "#999933",
  "Finding funding" = "#DDCC77",
  "Managing issues" = "#CC6677",
  "Attracting users" = "#882255"
)
```

```

top5_by_perc_plot <- ggplot(
  filtered_props,
  aes(
    x = job_category,
    y = proportion,
    fill = challenge
  )
) +
  geom_col(position = position_dodge()) +
  scale_y_continuous(labels = scales::percent, limits = c(0, 1)) +
  scale_fill_manual(values = chall_colors2) +
  labs(
    x = "Job Category",
    y = "Percent Saying\n'Frequently' or 'Always'",
    fill = "Challenge",
    title = "Top 5 Challenges by Job Category (Percent)"
  ) +
  theme(
    axis.title.x = element_blank(),
    axis.title.y = element_text(size = 24),
    axis.text.x = element_text(angle = 60, vjust = 0.6, size = 18),
    axis.text.y = element_text(size = 18),
    axis.ticks.x = element_blank(),
    legend.title = element_blank(),
    legend.text = element_text(size = 18),
    panel.background = element_blank(),
    panel.grid = element_line(linetype = "solid", color = "gray90"),
    plot.title = element_text(hjust = 0, size = 24, face = "bold"),
    plot.margin = unit(c(0.3, 0.3, 0.3, 0.3), "cm")
  )
)
top5_by_perc_plot

```

## Top 5 Challenges by Job Category (Percent)



```
save_plot("top5_challenges_by_job_percent.tiff", 12, 10, p=top5_by_perc_plot)
```

```
p_combined <- patchwork::wrap_plots(  
  top5_by_points_plot,  
  plot_spacer(),  
  top5_by_perc_plot  
) +  
  plot_layout(widths = c(1, 0.05, 1)) +  
  theme(plot.margin = margin(t = 1, r = 1, b = 1, l = 1, unit = "cm"))  
p_combined <- p_combined +  
  plot_annotation(tag_levels = "A") &  
  theme(plot.tag = element_text(size = 26))  
  
# SVG is higher quality  
svglite::svglite(  
  file.path(FIGURE_PATH, "figureS7.svg"),  
  width = 26,  
  height = 10  
)
```

```
print(p_combined)
dev.off()
```

```
pdf
2
```

## Consider clusters

### Exploratory plot

In a previous notebook, we found that the distributions of responses to the various challenges could be clustered like so:

Cluster 1:

Education time  
Documentation time  
Coding time

Cluster 2:

Securing funding  
Hiring  
Finding funding

Cluster 3:

Everything else

This makes me curious: does the distribution of job categories also vary by cluster? Before we try any statistics, let's just make a plot. This will be a variation of the detailed plot above.

We're just going to subset the "frequently" or "always" data to include only clusters 1 and 2, and we'll reorder the factor levels accordingly.

```
clusters1and2 <- c(
  "Education time",
  "Documentation time",
  "Coding time",
  "Securing funding",
  "Finding funding",
  "Hiring"
)
```

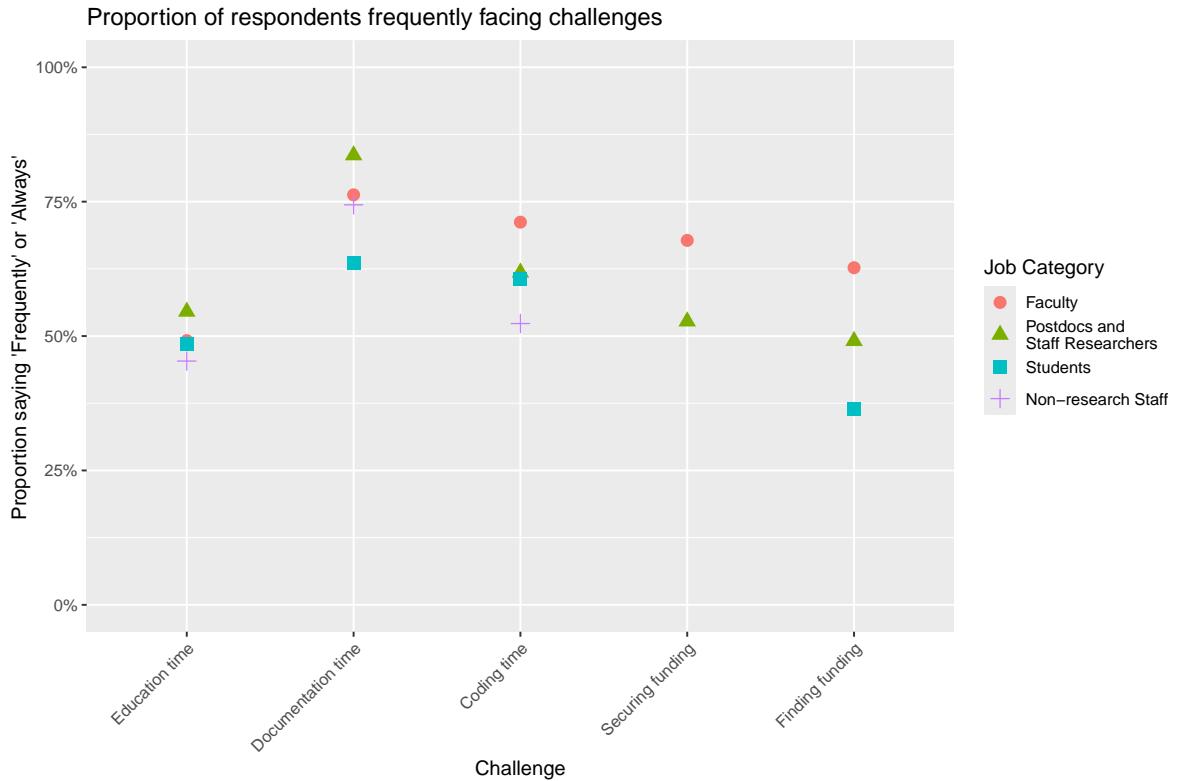
```

to_plot_clusters <- subset(filtered_props, challenge %in% clusters1and2)

to_plot_clusters$challenge <- factor(
  to_plot_clusters$challenge,
  levels = clusters1and2
)

challenges_plot_clusters1and2 <- ggplot(
  to_plot_clusters,
  aes(
    x = challenge,
    y = proportion,
    group = job_category,
    color = job_category,
    shape = job_category
  )
) +
  geom_point(size = 3) +
  scale_y_continuous(labels = scales::percent, limits = c(0, 1)) +
  labs(
    x = "Challenge",
    y = "Proportion saying 'Frequently' or 'Always'",
    color = "Job Category",
    shape = "Job Category",
    title = "Proportion of respondents frequently facing challenges"
  ) +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
challenges_plot_clusters1and2

```



Hm. Well, the results for cluster 1 are a little messy, which kind of makes sense, since you'd expect these to be challenges everyone struggles with. The only obvious trend to me is that undergrads struggle less than everyone else with education time and documentation time. But this group is too small to conclude anything with confidence.

Cluster 2 is a bit more interesting. It seems, at a glance, like it's pretty safe to say that faculty struggle with these challenges more than everyone else, with postdocs and staff researchers close behind.

I guess we should do a regression to test it.

## Regression for cluster 2

Let's once again combine the smaller groups to get more statistical power.

```
cluster2data <- subset(
  long_data,
  challenge %in% c("Securing funding", "Finding funding", "Hiring")
)
```

```

cluster2data <- cluster2data %>%
  mutate(
    job_category = recode(
      job_category,
      "Post-Doc" = "Postdocs and Staff Researchers",
      "Other research staff" = "Postdocs and Staff Researchers"
    )
  )

cluster2data <- cluster2data %>%
  mutate(
    job_category = recode(
      job_category,
      "Grad Student" = "Students",
      "Undergraduate" = "Students"
    )
  )

cluster2data$challenge_level <- factor(cluster2data$challenge_level)

cluster2data

```

	job_category	participantID	challenge	challenge_level	challenge_score
	<chr>	<chr>	<chr>	<fct>	<int>
1	Faculty	1	Hiring	Always	4
2	Faculty	1	Finding	~ Always	4
3	Faculty	1	Securing	~ Always	4
4	Postdocs and Staff R~	2	Hiring	Rarely	1
5	Postdocs and Staff R~	2	Finding	~ Frequently	3
6	Postdocs and Staff R~	2	Securing	~ Occasionally	2
7	Postdocs and Staff R~	3	Hiring	Frequently	3
8	Postdocs and Staff R~	3	Finding	~ Always	4
9	Postdocs and Staff R~	3	Securing	~ Always	4
10	Faculty	4	Hiring	Always	4
# i 689 more rows					

This code is really similar to code in solutions\_stats.qmd. See that notebook for commentary on these models.

### **Model 1: job\_category \* challenge interaction**

```
fit1 <- ordinal::clmm(challenge_level ~ job_category * challenge +  
  (1 | participantID),  
  data = cluster2data, link = "logit", Hess = TRUE)
```

### **Model 2: challenge as a random effect, no correlation between participant intercept and job effect**

```
fit2 <- ordinal::clmm(challenge_level ~ job_category +  
  (1 | challenge) +  
  (1 | participantID) +  
  (0 + job_category | challenge),  
  data = cluster2data, link = "logit", Hess = TRUE)
```

### **Model 3: No job category**

```
fit3 <- ordinal::clmm(challenge_level ~ challenge +  
  (1 | participantID),  
  data = cluster2data, link = "logit", Hess = TRUE)
```

### **Model 4: No challenge category**

```
fit4 <- ordinal::clmm(challenge_level ~ job_category +  
  (1 | participantID),  
  data = cluster2data, link = "logit", Hess = TRUE)
```

### **Model 5: job\_category + solution**

```
fit5 <- ordinal::clmm(challenge_level ~ job_category + challenge +  
  (1 | participantID),  
  data = cluster2data, link = "logit", Hess = TRUE)
```

## Model 6: no random effects

```
# note clm function bc clmm is for mixed models
fit6 <- ordinal::clm(challenge_level ~ job_category * challenge,
                      data = cluster2data, link = "logit", Hess = TRUE)
```

## Compare models

```
models <- list(
  "fit1"=fit1, # job_category * challenge
  "fit2"=fit2, # challenge as random effect
  "fit3"=fit3, # Null model: no job
  "fit4"=fit4, # Null model: no challenge
  "fit5"=fit5, # Null model: no interaction
  "fit6"=fit6 # Null model: no participants
)
sapply(models, function(x) round(stats::AIC(x)))
```

```
fit1 fit2 fit3 fit4 fit5 fit6
2096 2112 2102 2107 2087 2188
```

Models 1 and 5 look best in terms of AIC.

Let's check the condition number of the Hessian. I don't really understand what this is, but the clmm2 tutorial says that high numbers, say larger than say  $10^4$  or  $10^6$ , indicate poor fit.

```
sapply(models, function(x)
  summary(x)$info["cond.H"]
)
```

```
$fit1.cond.H
[1] "9.0e+02"

$fit2.cond.H
[1] "5.2e+02"
```

```

$fit3.cond.H
[1] "1.0e+02"

$fit4.cond.H
[1] "3.6e+02"

$fit5.cond.H
[1] "3.7e+02"

$fit6.cond.H
[1] "1.5e+03"

```

All look ok.

### Complex models vs null models

Let's use an anova to compare nested models.

```
stats::anova(fit1, fit5)
```

Likelihood ratio tests of cumulative link models:

```

formula:                                         link:
fit5 challenge_level ~ job_category + challenge + (1 | participantID) logit
fit1 challenge_level ~ job_category * challenge + (1 | participantID) logit
threshold:
fit5 flexible
fit1 flexible

no.par      AIC  logLik LR.stat df Pr(>Chisq)
fit5      11 2086.8 -1032.4
fit1      17 2095.9 -1031.0  2.8688   6     0.8251

```

Interesting, that p-value is not significant. So it appears the interaction term is not needed.

Let's also double-check that participants are worth including.

```
stats::anova(fit1, fit6)
```

```

Likelihood ratio tests of cumulative link models:

      formula:                                         link:
fit6 challenge_level ~ job_category * challenge      logit
fit1 challenge_level ~ job_category * challenge + (1 | participantID) logit
      threshold:
fit6 flexible
fit1 flexible

      no.par    AIC  logLik LR.stat df Pr(>Chisq)
fit6     16 2188.2 -1078.1
fit1     17 2095.9 -1031.0  94.252  1 < 2.2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Yes, it appears they are.

Does it matter whether we include job as a variable? Let's compare it to the model without an interaction term.

```
stats::anova(fit3, fit5)
```

```
Likelihood ratio tests of cumulative link models:
```

```

      formula:                                         link:
fit3 challenge_level ~ challenge + (1 | participantID)      logit
fit5 challenge_level ~ job_category + challenge + (1 | participantID) logit
      threshold:
fit3 flexible
fit5 flexible

      no.par    AIC  logLik LR.stat df Pr(>Chisq)
fit3     8 2102.0 -1043.0
fit5    11 2086.8 -1032.4  21.189  3  9.616e-05 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Yes, including job improves model fit.

So far, fit5 is the one to beat.

## More goodness-of-fit tests

SEs of the coefficients

```
summary(fit5$coefficients)
```

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
	-0.7551	0.5609	0.8817	1.3455	1.4904	4.8889

```
summary(fit2$coefficients)
```

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
	-0.9148	0.6869	1.1423	1.5189	1.9698	4.7287

These look pretty similar.

Let's do one more diagnostic. fit6 is the equivalent model to fit1b, but with fixed effects only. Since we can do the nominal\_test and scale\_test on this model, let's try it and see if it sets off any red flags.

```
nominal_test(fit6)
```

Tests of nominal effects

```
formula: challenge_level ~ job_category * challenge
          Df logLik    AIC     LRT Pr(>Chi)
<none>           -1078.1 2188.2
job_category      12 -1058.8 2173.5 38.635 0.0001208 ***
challenge         8 -1070.4 2188.8 15.345 0.0527751 .
job_category:challenge
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
scale_test(fit6)
```

Tests of scale effects

```
formula: challenge_level ~ job_category * challenge
          Df logLik    AIC     LRT Pr(>Chi)
<none>           -1078.1 2188.2
```

```

job_category           3 -1069.8 2177.6 16.552 0.0008736 ***
challenge              2 -1072.3 2180.7 11.463 0.0032428 **
job_category:challenge 11 -1059.4 2172.8 37.383 9.939e-05 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Boo. The model with no random effects has violations to both assumptions.

Ouch. That's not ideal. Maybe we can proceed with caution, and follow up with a non-parametric test on whatever trends we find? <https://www.fharrell.com/post/po/>

## EMMs

Again, see the solutions\_stats notebook for more detail on this.

```

emm <- summary(emmeans(fit5, ~ challenge | job_category, mode = "mean.class"))
emm

job_category = Faculty:
challenge      mean.class    SE  df asymp.LCL asymp.UCL
Finding funding     2.36 0.194 Inf     1.98     2.74
Hiring            2.84 0.203 Inf     2.44     3.23
Securing funding   2.24 0.188 Inf     1.87     2.61

job_category = Non-research Staff:
challenge      mean.class    SE  df asymp.LCL asymp.UCL
Finding funding     3.39 0.152 Inf     3.09     3.69
Hiring             3.80 0.137 Inf     3.54     4.07
Securing funding   3.27 0.157 Inf     2.96     3.58

job_category = Postdocs and Staff Researchers:
challenge      mean.class    SE  df asymp.LCL asymp.UCL
Finding funding     2.93 0.204 Inf     2.53     3.33
Hiring             3.39 0.186 Inf     3.02     3.75
Securing funding   2.80 0.205 Inf     2.40     3.20

job_category = Students:
challenge      mean.class    SE  df asymp.LCL asymp.UCL
Finding funding     3.40 0.230 Inf     2.95     3.85
Hiring             3.81 0.205 Inf     3.41     4.21
Securing funding   3.28 0.237 Inf     2.82     3.75

```

```
Confidence level used: 0.95
```

## Pairwise comparisons and p-values

Here we look at pairwise contrasts by challenge.

```
emm2 <- emmeans(fit5, ~ job_category | challenge, mode = "mean.class")
by_chall <- summary(
  pairs(emm2, by = "challenge"),
  infer = TRUE # infer CIs
)

by_chall

challenge = Finding funding:
contrast                               estimate   SE  df
Faculty - (Non-research Staff)          -1.03047 0.227 Inf
Faculty - Postdocs and Staff Researchers -0.56513 0.261 Inf
Faculty - Students                      -1.03984 0.284 Inf
(Non-research Staff) - Postdocs and Staff Researchers  0.46535 0.233 Inf
(Non-research Staff) - Students          -0.00936 0.258 Inf
Postdocs and Staff Researchers - Students -0.47471 0.289 Inf
asymp.LCL asymp.UCL z.ratio p.value
  -1.613    -0.448   -4.542  <.0001
  -1.237     0.107   -2.162  0.1341
  -1.770    -0.309   -3.657  0.0015
  -0.134     1.065   1.994  0.1900
  -0.673     0.654   -0.036  1.0000
  -1.218     0.269   -1.640  0.3560

challenge = Hiring:
contrast                               estimate   SE  df
Faculty - (Non-research Staff)          -0.96846 0.225 Inf
Faculty - Postdocs and Staff Researchers -0.55074 0.256 Inf
Faculty - Students                      -0.97680 0.272 Inf
(Non-research Staff) - Postdocs and Staff Researchers  0.41772 0.212 Inf
(Non-research Staff) - Students          -0.00834 0.230 Inf
Postdocs and Staff Researchers - Students -0.42606 0.261 Inf
asymp.LCL asymp.UCL z.ratio p.value
  -1.546    -0.391   -4.310  0.0001
```

```

-1.208      0.107  -2.151   0.1371
-1.675      -0.279  -3.596   0.0018
-0.128      0.963   1.968   0.2002
-0.599      0.582   -0.036   1.0000
-1.098      0.246   -1.630   0.3617

challenge = Securing funding:
contrast                                         estimate    SE  df
Faculty - (Non-research Staff)                 -1.03296 0.225 Inf
Faculty - Postdocs and Staff Researchers       -0.55827 0.259 Inf
Faculty - Students                            -1.04261 0.286 Inf
(Non-research Staff) - Postdocs and Staff Researchers 0.47469 0.237 Inf
(Non-research Staff) - Students                -0.00965 0.266 Inf
Postdocs and Staff Researchers - Students      -0.48434 0.295 Inf
asympt.LCL asympt.UCL z.ratio p.value
  -1.611     -0.455  -4.589 <.0001
  -1.223     0.106   -2.159  0.1349
  -1.778     -0.307  -3.643  0.0015
  -0.134     1.083   2.005  0.1862
  -0.693     0.674   -0.036  1.0000
  -1.243     0.275   -1.640  0.3563

```

Confidence level used: 0.95

Conf-level adjustment: tukey method for comparing a family of 4 estimates

P value adjustment: tukey method for comparing a family of 4 estimates

Wow, the p-values are really similar across the board. Faculty rate these challenges higher than students and NR staff, but not higher than postdocs and staff researchers.

## Kruskal-Wallis test for ranking differences between groups

Non-parametric test for the extent of disagreement between groups. Whereas above, we tested for differences in mean ratings, here we are testing for differences in the distributions of ratings for each solution.

```

cluster2data_numcoded <- cluster2data %>%
  mutate(
    challenge_score = recode(
      challenge_level,
      "Non-applicable" = 0L,
      "Never" = 0L,

```

```

    "Rarely" = 1L,
    "Occasionally" = 2L,
    "Frequently" = 3L,
    "Always" = 4L
  )
)

cluster2data_numcoded$job_category <- factor(cluster2data_numcoded$job_category)

kw_results <- sapply(split(cluster2data_numcoded, cluster2data_numcoded$challenge), function
  kruskal.test(challenge_score ~ job_category, data = df)$p.value
)

p_adj_kw <- p.adjust(kw_results, "BH")

p_adj_kw < 0.05

```

Finding funding	Hiring	Securing funding
TRUE	TRUE	TRUE

```
sum(p_adj_kw < 0.05)
```

```
[1] 3
```

Ok, so K-W test indicates that for all three challenges, there are differences between the groups.

Dunn test as a post-hoc test to see which groups are different from each other.

```

pairwise_results <- lapply(unique(cluster2data_numcoded$challenge), function(chall) {
  df <- subset(cluster2data_numcoded, challenge == chall)
  out <- FSA::dunnTest(challenge_score ~ job_category, data = df, method = "bh")
  cbind(challenge = chall, out$res)
})
pairwise_results <- do.call(rbind, pairwise_results)

```

Let's print the significant pairs.

```
subset(pairwise_results, P.adj < 0.05)
```

	challenge	Comparison	
1	Hiring	Faculty - Non-research Staff	
2	Hiring	Faculty - Postdocs and Staff Researchers	
4	Hiring	Faculty - Students	
7	Funding funding	Faculty - Non-research Staff	
9	Funding funding	Non-research Staff - Postdocs and Staff Researchers	
10	Funding funding	Faculty - Students	
13	Securing funding	Faculty - Non-research Staff	
15	Securing funding	Non-research Staff - Postdocs and Staff Researchers	
16	Securing funding	Faculty - Students	
	Z P.unadj P.adj		
1	3.202369 1.363021e-03 4.089062e-03		
2	2.483752 1.300062e-02 2.600123e-02		
4	4.191500 2.771168e-05 1.662701e-04		
7	4.319614 1.563026e-05 9.378155e-05		
9	-2.899242 3.740657e-03 1.122197e-02		
10	2.664241 7.716231e-03 1.543246e-02		
13	4.741135 2.125239e-06 1.275143e-05		
15	-3.077589 2.086828e-03 4.173655e-03		
16	3.174243 1.502280e-03 4.506841e-03		

And the non-significant pairs

```
subset(pairwise_results, P.adj >= 0.05)
```

	challenge	Comparison	
3	Hiring	Non-research Staff - Postdocs and Staff Researchers	
5	Hiring	Non-research Staff - Students	
6	Hiring	Postdocs and Staff Researchers - Students	
8	Funding funding	Faculty - Postdocs and Staff Researchers	
11	Funding funding	Non-research Staff - Students	
12	Funding funding	Postdocs and Staff Researchers - Students	
14	Securing funding	Faculty - Postdocs and Staff Researchers	
17	Securing funding	Non-research Staff - Students	
18	Securing funding	Postdocs and Staff Researchers - Students	
	Z P.unadj P.adj		
3	-0.4391213 0.66057367 0.66057367		
5	1.8058145 0.07094732 0.08513678		
6	2.0236538 0.04300579 0.06450868		
8	1.2252483 0.22048168 0.26457802		
11	-0.7377959 0.46063850 0.46063850		
12	1.5871988 0.11246762 0.16870143		

```
14 1.4411382 0.14954564 0.17945476  
17 -0.5443851 0.58617651 0.58617651  
18 1.9069060 0.05653276 0.08479915
```

Cool. In all three cases, faculty are significantly different from NR staff and students.

```
sessionInfo()
```

```
R version 4.4.2 (2024-10-31)  
Platform: aarch64-apple-darwin20  
Running under: macOS 26.1  
  
Matrix products: default  
BLAS: /Library/Frameworks/R.framework/Versions/4.4-arm64/Resources/lib/libRblas.0.dylib  
LAPACK: /Library/Frameworks/R.framework/Versions/4.4-arm64/Resources/lib/libRlapack.dylib;  
  
locale:  
[1] C.UTF-8/C.UTF-8/C.UTF-8/C.UTF-8/C.UTF-8/C.UTF-8  
  
time zone: America/Los_Angeles  
tzcode source: internal  
  
attached base packages:  
[1] tools      grid       stats      graphics   grDevices datasets  utils  
[8] methods     base  
  
other attached packages:  
[1] treemapify_2.5.6    tidyverse_1.3.1      svglite_2.2.1  
[4] stringr_1.5.1      scales_1.4.0        readr_2.1.5  
[7] pwr_1.3-0          patchwork_1.3.2    ordinal_2023.12-4.1  
[10] lme4_1.1-37        Matrix_1.7-1       languageserver_0.3.16  
[13] here_1.0.1         gtools_3.9.5       ggforce_0.5.0  
[16] FSA_0.10.0         fpc_2.2-13       forcats_1.0.0  
[19] factoextra_1.0.7   ggplot2_3.5.2     emmeans_1.11.2  
[22] dplyr_1.1.4        corrplot_0.95     ComplexHeatmap_2.22.0  
[25] cluster_2.1.8.1    BiocManager_1.30.26  
  
loaded via a namespace (and not attached):  
[1] Rdpack_2.6.4        dunn.test_1.3.6     rlang_1.1.6  
[4] magrittr_2.0.3       clue_0.3-66        GetoptLong_1.0.5  
[7] matrixStats_1.5.0    compiler_4.4.2     flexmix_2.3-20  
[10] systemfonts_1.2.3   png_0.1-8         callr_3.7.6
```

[13]	vctrs_0.6.5	pkgconfig_2.0.3	shape_1.4.6.1
[16]	crayon_1.5.3	fastmap_1.2.0	labeling_0.4.3
[19]	utf8_1.2.6	rmarkdown_2.29	ggfittext_0.10.2
[22]	tzdb_0.5.0	ps_1.9.1	nloptr_2.2.1
[25]	purrrr_1.1.0	xfun_0.53	modeltools_0.2-24
[28]	jsonlite_2.0.0	tweenr_2.0.3	parallel_4.4.2
[31]	prabclus_2.3-4	R6_2.6.1	stringi_1.8.7
[34]	RColorBrewer_1.1-3	boot_1.3-31	diptest_0.77-2
[37]	numDeriv_2016.8-1.1	estimability_1.5.1	Rcpp_1.1.0
[40]	iterators_1.0.14	knitr_1.50	IRanges_2.40.1
[43]	splines_4.4.2	nnet_7.3-19	tidyselect_1.2.1
[46]	yaml_2.3.10	doParallel_1.0.17	codetools_0.2-20
[49]	processx_3.8.6	lattice_0.22-6	tibble_3.3.0
[52]	withr_3.0.2	evaluate_1.0.4	polyclip_1.10-7
[55]	xml2_1.4.0	circlize_0.4.16	mclust_6.1.1
[58]	kernlab_0.9-33	pillar_1.11.0	renv_1.1.5
[61]	foreach_1.5.2	stats4_4.4.2	reformulas_0.4.1
[64]	generics_0.1.4	rprojroot_2.1.1	S4Vectors_0.44.0
[67]	hms_1.1.3	minqa_1.2.8	xtable_1.8-4
[70]	class_7.3-22	glue_1.8.0	robustbase_0.99-4-1
[73]	mvtnorm_1.3-3	rbibutils_2.3	colorspace_2.1-1
[76]	nlme_3.1-166	cli_3.6.5	textshaping_1.0.1
[79]	gttable_0.3.6	DEoptimR_1.1-4	digest_0.6.37
[82]	BiocGenerics_0.52.0	ucminf_1.2.2	ggrepel_0.9.6
[85]	rjson_0.2.23	farver_2.1.2	htmltools_0.5.8.1
[88]	lifecycle_1.0.4	GlobalOptions_0.1.2	MASS_7.3-61