

Contributor roles

Overview

This script explores Q4: “Which of these open source contributor roles has ever applied to you?”.

Import packages and utilities

```
project_root <- here::here() # requires that you be somewhere in the
# project directory (not above it)
# packages
suppressMessages(source(file.path(project_root, "scripts/packages.R")))
# functions and objects used across scripts
suppressMessages(source(file.path(project_root, "scripts/utils.R")))
```

Load data

```
roles_raw <- load_qualtrics_data("clean_data/contributor_roles_Q4.tsv")
other_quant <- load_qualtrics_data("clean_data/other_quant.tsv")
head(roles_raw)
```

	Maintainer	Contributor	Bug/Issue Reporter	Community Manager	Educator	Other
1	1	1	1	1	1	0
2	0	1	0	0	0	0
3	1	1	1	0	1	0
4	1	1	1	0	1	0
5	1	1	1	0	1	0

6	0	0	0	0	0	0
	Supervisor	IT/Systems administrator	UI/UX Designer	Technical support		
1	1		0	0		1
2	0		0	0		0
3	0		0	1		0
4	1		0	0		0
5	1		0	0		0
6	0		0	0		0

```
nrow(roles_raw)
```

```
[1] 332
```

Wrangle data

Drop rows with all 0s (NAs were already converted to 0s during data cleanup).

```
roles <- filter_all(roles_raw, any_vars(. != 0))
nrow(roles)
```

```
[1] 233
```

Explore data

```
counts <- colSums(roles)
counts <- data.frame(Role = names(counts), Count = as.integer(counts), row.names = NULL)
counts <- counts %>%
  arrange(desc(Count))

counts
```

	Role	Count
1	Bug/Issue Reporter	189
2	Contributor	187
3	Educator	138
4	Maintainer	134
5	Technical support	102

6	Supervisor	96
7	Community Manager	49
8	UI/UX Designer	46
9	IT/Systems administrator	44
10	Other	13

That's mildly interesting. There's about 2.7-3 maintainers for every community manager, sys admin, or UI designer (keeping in mind that people might be both). I'm curious how many people are EXCLUSIVELY one thing. Mainly interested in how many people are EXCLUSIVELY a maintainer, contributor, bug reporter, or supervisor.

```
indices <- which(rowSums(roles) == 1 )
exclusive <- roles[indices,]
nrow(exclusive)
```

```
[1] 16
```

Let's repeat what we did for the whole data frame, but now with just people who filled one role.

```
counts_exc <- colSums(exclusive)
counts_exc <- data.frame(Role = names(counts_exc), Count = as.integer(counts_exc), row.names = names(counts_exc))
counts_exc <- counts_exc %>%
  arrange(desc(Count))

counts_exc
```

	Role	Count
1	Bug/Issue Reporter	5
2	Contributor	4
3	Supervisor	3
4	Educator	2
5	Maintainer	1
6	Community Manager	1
7	Other	0
8	IT/Systems administrator	0
9	UI/UX Designer	0
10	Technical support	0

So, out of 233 contributors, only 16 identified with exactly one contributor role. 5 of those were Bug/Issue Reporter, the most common exclusive-role.

```

# Calculate total roles for each participant
roles2 <- cbind(roles, total_roles = rowSums(roles))

# Get mean number of roles per column (role)
get_mean_num_roles <- function(df, colnum) {
  # drop rows where the entry in this col is 0
  filtered <- df[df[[colnum]] != 0, ]
  return(
    mean(filtered$total_roles)
  )
}
num_roles <- data.frame(
  role = names(roles),
  avg_num_roles_per_participant = sapply(
    seq(ncol(roles)),
    function(x) get_mean_num_roles(roles2, x)
  )
)

num_roles <- num_roles[
  order(
    num_roles$avg_num_roles_per_participant,
    decreasing = TRUE
  ),
]

num_roles

```

	role	avg_num_roles_per_participant
8	IT/Systems administrator	6.386364
6	Other	6.384615
9	UI/UX Designer	6.369565
4	Community Manager	6.163265
10	Technical support	5.715686
7	Supervisor	5.583333
5	Educator	5.195652
1	Maintainer	5.194030
2	Contributor	4.737968
3	Bug/Issue Reporter	4.682540

Huh. So I guess no one role is particularly enriched for people who hold many roles. (I mean, we can't really get statistical significance with just one observation, and this doesn't feel worth

modeling.) I guess bug reporters and contributors have the fewest roles, by a small margin, which makes sense.

Three-way comparison of maintainers, contributors, and bug reporters

I'd be interested in the 3-way venn diagram of contributors, but reporters and maintainers. (Though, if we do plot this, an UpSet plot would probably be better than a Venn diagram)

```
# maintainers, contributors, and bug reporters in that order
three_way_counts <- with(roles, {
  M <- Maintainer == 1
  C <- Contributor == 1
  B <- `Bug/Issue Reporter` == 1

  c(
    "M, !C, !B" = sum(M & !C & !B),
    "M, C, !B" = sum(M & C & !B),
    "M, !C, B" = sum(M & !C & B),
    "M, C, B" = sum(M & C & B),
    "!M, !C, B" = sum(!M & !C & B),
    "!M, C, B" = sum(!M & C & B),
    "!M, C, !B" = sum(!M & C & !B),
    "!M, !C, !B" = sum(!M & !C & !B)
  )
})

three_way_counts
```

M, !C, !B	M, C, !B	M, !C, B	M, C, B	!M, !C, B	!M, C, B	!M, C, !B
5	14	4	111	27	47	15
!M, !C, !B						
10						

```
# This should be 233, the number of experienced contributors
sum(three_way_counts)
```

```
[1] 233
```

```
# Number of people who contributed as at least one of these three roles
at_least_one <- sum(three_way_counts)-three_way_counts["!M, !C, !B"]
at_least_one
```

```
!M, !C, !B
      223
```

```
# Percent of total who identified as at least one of the three
round(three_way_counts/at_least_one*100, digits=1)
```

```

M, !C, !B   M, C, !B   M, !C, B   M, C, B   !M, !C, B   !M, C, B   !M, C, !B
      2.2      6.3      1.8      49.8      12.1      21.1      6.7
!M, !C, !B
      4.5
```

Interesting. So 96% of the survey respondents identified as at least one of Maintainer, Contributor, or Bug Reporter (223/233). 50.0% of those folks identified as all three (111/223). 21.1% identified as a contributor and bug reporter, but not a maintainer, and 12.1% identified as a bug reporter only. Together, these three groups make up 50+21.1+12.2 83.3% of the total.

UpSet plot

```
# turn named vector into a table of membership + counts
comb_tbl <- tibble::enframe(three_way_counts, name = "pattern", value = "n") %>%
  mutate(
    M = !str_detect(pattern, "!M"),
    C = !str_detect(pattern, "!C"),
    B = !str_detect(pattern, "!B")
  ) %>%
  select(M, C, B, n)

comb_tbl
```

```
# A tibble: 8 x 4
  M     C     B     n
  <lgl> <lgl> <lgl> <int>
1 TRUE  FALSE FALSE     5
2 TRUE  TRUE  FALSE    14
```

3	TRUE	FALSE	TRUE	4
4	TRUE	TRUE	TRUE	111
5	FALSE	FALSE	TRUE	27
6	FALSE	TRUE	TRUE	47
7	FALSE	TRUE	FALSE	15
8	FALSE	FALSE	FALSE	10

```
# expand rows by counts (one row per element)
elem_tbl <- comb_tbl %>% uncount(n)
```

```
names(elem_tbl) <- c(
  "Maintainer",
  "Contributor",
  "Bug reporter"
)
elem_tbl
```

```
# A tibble: 233 x 3
  Maintainer Contributor `Bug reporter`
  <lgl>         <lgl>         <lgl>
1 TRUE         FALSE         FALSE
2 TRUE         FALSE         FALSE
3 TRUE         FALSE         FALSE
4 TRUE         FALSE         FALSE
5 TRUE         FALSE         FALSE
6 TRUE         TRUE          FALSE
7 TRUE         TRUE          FALSE
8 TRUE         TRUE          FALSE
9 TRUE         TRUE          FALSE
10 TRUE        TRUE          FALSE
# i 223 more rows
```

```
m <- make_comb_mat(
  elem_tbl[, c("Maintainer", "Contributor", "Bug reporter")],
  mode = "distinct"
)
# distinct is the default and indicates that the combinations are
# mutually exclusive; each data point appears only once in the data set

ordered_combs <- order(comb_size(m), decreasing = TRUE)
```

```

# Since this function call has a lot of args, I'm breaking it up into chunks

# Control bars at top of plot
top_ann_args <- ComplexHeatmap::upset_top_annotation(
  m,
  add_numbers = TRUE,
  numbers_gp = gpar(fontsize = 14), # bigger labels above bars
  height = unit(70, "mm"), # bigger bars
  annotation_name_gp = gpar(fontsize = 14), # bigger axis label ("Intersection size")
  axis_param = list(gp = gpar(fontsize = 11)) # bigger axis tick labels
)

right_ann_args <- ComplexHeatmap::upset_right_annotation(
  m,
  width = unit(40, "mm"), # width of right bars
  annotation_name_gp = gpar(fontsize = 14), # bigger axis label ("Set size")
  axis_param = list(gp = gpar(fontsize = 11)) # bigger axis tick labels
)


upset_plot <- UpSet(
  m,
  # These args control the membership block
  height = unit(32, "mm"), # block height
  pt_size = unit(3, "mm"), # dot size in the membership matrix
  row_names_gp = gpar(fontsize = 14), # row labels (set names)

  comb_order = ordered_combs,
  set_order = c("Maintainer", "Contributor", "Bug reporter"), # control set order

  top_annotation = top_ann_args,
  right_annotation = right_ann_args,
  column_title = "Participants' open source roles",
  column_title_gp = gpar(fontsize = 22, fontface = "bold")
)

upset_plot

```

Save the plot

```
tiff(
    file.path(FIGURE_PATH, "roles_upset.tiff"),
    width = 10,
    height = 6,
    unit = "in",
    res = 700
)
upset_plot
dev.off()
```

pdf
2

Further exploration with maintainers

Out of curiosity, how many maintainers also selected tech support?

```
two_way_counts <- with(roles, {  
  M <- Maintainer == 1  
  TS <- `Technical support` == 1  
  
  c(  
    "M, !T" = sum(M & !TS),  
    "M, T" = sum(M & TS)  
  )  
})  
  
two_way_counts
```

```
M, !T M, T  
66 68
```

Eh, not that interesting. I thought more maintainers would also be tech supporters.

What's the ratio of maintainers to other contributors?

```
nrow(subset(roles, Maintainer == 1))
```

```
[1] 134
```

```
nrow(subset(roles, Maintainer == 0))
```

```
[1] 99
```

That's somewhat interesting, and I guess it's consistent with the "three-way venn-diagram" above. 134 people identified as maintainers, and only 99 people identified as contributors of some sort, but not maintainers. So $134/233 = 57.5\%$ of experienced contributors in this survey are maintainers. I'm getting the sense, basically, that there are a lot of maintainers in this survey.

Bring in field of study

Are Math/CS people more likely to be maintainers than biologists?

```
roles_field <- cbind(roles_raw, other_quant$field_of_study)
# Rename column
names(roles_field)[ncol(roles_field)] <- "field_of_study"
# Filter out non-academics (people who didn't answer the field of study question)
roles_field <- roles_field %>%
  filter(field_of_study != "")

nrow(roles_field)
```

[1] 188

```
# Remove people who did not select any role
roles_field_clean <- roles_field[rowSums(roles_field[, -ncol(roles_field)]) != 0,]

nrow(roles_field_clean)
```

[1] 147

```
# Sanity check: make sure none of the rows sums to 0
unnname(rowSums(roles_field_clean %>% select(-field_of_study)))
```

```
[1] 7 1 5 5 5 4 5 3 6 4 3 2 3 3 5 2 3 4 2 4 8 3 5 3 3
[26] 4 1 4 2 3 3 3 8 4 2 3 6 1 8 10 4 5 2 3 5 4 3 3 3 4
[51] 4 7 2 6 6 5 4 2 7 7 3 4 3 5 2 2 2 5 9 5 6 2 7 6 1
[76] 2 8 2 5 4 7 4 3 4 3 3 4 2 5 1 2 3 3 4 4 4 10 4 6 4
[101] 4 6 5 4 3 1 1 2 6 7 1 1 7 8 5 4 2 4 6 2 6 5 8 4 8
[126] 7 4 7 7 3 1 2 6 4 4 1 4 5 3 5 6 2 8 4 3 6 2
```

Okay, so we have a total of 188 academics. It looks like 147 of those are experienced contributors—people who selected at least one role.

```
maintainer_props <- roles_field_clean %>%
  group_by(field_of_study) %>%
  summarise(
    n_people = n(),
```

```

    n_maintainers      = sum(Maintainer == 1),
    prop_maintainers   = mean(Maintainer == 1)
  ) %>%
  arrange(desc(prop_maintainers))

```

```
maintainer_props
```

```

# A tibble: 5 x 4
  field_of_study    n_people n_maintainers prop_maintainers
  <chr>            <int>      <int>          <dbl>
1 Math and CS      72         53           0.736
2 Physical sciences 27         18           0.667
3 Life sciences    34         18           0.529
4 Social sciences  10          4           0.4
5 Humanities        4          1           0.25

```

Meh, mildly interesting. 50% or more of our survey respondents from STEM fields are maintainers. There's really not sufficient respondents from the non-STEM fields to draw any conclusions there.

Campus

What proportion of respondents are maintainers, per campus?

```

roles_campus <- cbind(roles_raw, other_quant$campus)
# Rename column
names(roles_campus)[ncol(roles_campus)] <- "campus"
# Filter out non-UC
roles_campus <- roles_campus %>%
  filter(campus != "I'm not affiliated with UC")
# Remove people who did not select any role
roles_campus <- roles_campus[rowSums(roles_campus[, -ncol(roles_campus)]) != 0,]

nrow(roles_campus)

```

```
[1] 233
```

```

maint_by_campus <- roles_campus %>%
  group_by(campus) %>%
  summarise(
    n_rows = n(),
    n_maintainer = sum(Maintainer == 1),
    prop_maintainer = mean(Maintainer == 1)
  ) %>%
  arrange(desc(prop_maintainer))

maint_by_campus <- maint_by_campus[
  order(
    maint_by_campus$prop_maintainer,
    decreasing = TRUE
  ),
]
maint_by_campus

```

```

# A tibble: 10 x 4
  campus      n_rows n_maintainer prop_maintainer
  <chr>      <int>      <int>         <dbl>
1 Other UC          19           15         0.789
2 UC Santa Cruz      32           22         0.688
3 UC Berkeley        26           16         0.615
4 UC San Francisco    7            4         0.571
5 UC Santa Barbara   61           34         0.557
6 UC Los Angeles     40           21         0.525
7 UC Merced           8            4          0.5
8 UC Davis           29           14         0.483
9 UC San Diego        9            4         0.444
10 UC Irvine          2            0          0

```

Meh, again, only mildly interesting. I suppose it's an interesting add-on to the stat that about half of our respondents are maintainers. Almost like error basr, where the variability comes from the campuses. The campuses with fewer than ten respondents just feel like noise to me, since we can't really draw conclusions. Let's filter those out.

```

subset(maint_by_campus, n_rows > 10)

```

```

# A tibble: 6 x 4
  campus      n_rows n_maintainer prop_maintainer

```

	<chr>	<int>	<int>	<dbl>
1	Other UC	19	15	0.789
2	UC Santa Cruz	32	22	0.688
3	UC Berkeley	26	16	0.615
4	UC Santa Barbara	61	34	0.557
5	UC Los Angeles	40	21	0.525
6	UC Davis	29	14	0.483

That's a little clearer.